

Economics Letters

Can Time-Series Foundation Models Generate Improved Multivariate Forecasts? --Manuscript Draft--

Manuscript Number:	EL67834
Article Type:	Short Communication
Keywords:	multivariate forecasting; foundation models, generative AI
Corresponding Author:	Sanjiv R Das, PhD Santa Clara University San Jose, California UNITED STATES OF AMERICA
First Author:	Sanjiv R Das, PhD
Order of Authors:	Sanjiv R Das, PhD
	Tarang Goyal, MS
	Mohini Yadav, MS
Abstract:	How capable are time-series foundation models? Using Chronos-2, an open-source time-series foundation model, we evaluate pretrained time-series models for economic and financial forecasting with an emphasis on whether multivariate (MV) inputs improve accuracy relative to univariate (UV) baselines. The study covers two panels---the Magnificent-7 equities and U.S. Treasury interest rates---as well as a combined panel, using rolling monthly evaluations from 2000--2025. We vary input window lengths and forecast horizons and report RMSE and MAPE. Across datasets, MV forecasts consistently outperform UV forecasts, with especially strong gains for interest rates and meaningful improvements for equities. Series-level comparisons show MV improvements in every case, and error dispersion is generally lower under MV inputs. We also provide parameter-heatmap and time-series visualizations. However, mixing time series across equity and interest rate markets reduces forecast accuracy, indicating that adding noisy context degrades model performance. Overall, the results indicate that foundation models can leverage cross-series information to improve forecast accuracy in finance, and that the benefits are strongest when related series are modeled jointly under disciplined rolling protocols.



April 29, 2026.

To: The Editor
Economics Letters

Dear Editor:

I am submitting the paper, titled "*Can Time-Series Foundation Models Generate Improved Multivariate Forecasts*" for publication consideration. The paper is co-authored with Tarang Goyal and Mohini Yadav.

The study evaluates the Chronos-2 open-source foundation model for economic and financial forecasting, specifically comparing multivariate (MV) and univariate (UV) approaches across 25 years of data (2000–2025). Just as LLMs generate sequences of words, Chronos-2 is trained on thousands of time series, and generates forward sequences of data based on multivariate histories. The main findings of the paper are:

- *Time-Series foundation models (TSFMs) mimic LLMs successfully:* Just as LLMs generate sequences of words, The Chronos TSFM is trained on thousands of time series, and generates forward sequences of data based on multivariate histories.
- *Multivariate Superiority:* Unlike LLMs, TSFMs can ingest multiple sequences for multivariate forecasting, showing higher accuracy and lower error dispersion across U.S. Treasury rates and "Magnificent-7" equities.
- *Accuracy:* Gains are most significant when using related series under disciplined rolling protocols.
- *Context Matters:* While joint modeling of related stock and interest rate series separately improved performance, mixing unrelated data (equities and interest rates) actually degraded accuracy due to "noisy context."
- *AI Methodology:* In addition to its financial findings, the paper serves as a demonstration of how AI-driven workflows and foundation models can be integrated into modern financial research.

A separate Appendix is provided articulating how AI was used in the production of this paper. We hope this will also be helpful to readers. The paper's word count, excluding tables, figures, and references, is 1928 (computed using the TeXcount package (<https://app.uio.no/ifi/texcount/>)).

Thank you for your time and consideration.

Sincerely,

Sanjiv R. Das | William and Janice Terry Professor of Finance and Data Science
Leavey School of Business | Santa Clara University
Tel: (408)-554-2776 | Email: srdas@scu.edu
<http://srdas.github.io/>

Can Time-Series Foundation Models Generate Improved Multivariate Forecasts?

Highlights

- *Time-Series foundation models (TSFMs) mimic LLMs successfully:* Just as LLMs generate sequences of words, The Chronos TSFM is trained on thousands of time series, and generates forward sequences of data based on multivariate histories.
- *Multivariate Superiority:* Unlike LLMs, TSFMs can ingest multiple sequences for multivariate forecasting, showing higher accuracy and lower error dispersion across U.S. Treasury rates and "Magnificent-7" equities.
- *Accuracy:* Gains are most significant when using related series under disciplined rolling protocols.
- *Context Matters:* While joint modeling of related stock and interest rate series separately improved performance, mixing unrelated data (equities and interest rates) actually degraded accuracy due to "noisy context."
- *AI Methodology:* In addition to its financial findings, the paper serves as a demonstration of how AI-driven workflows and foundation models can be integrated into modern financial research.

Can Time-Series Foundation Models Generate Improved Multivariate Forecasts?

Authors:

- Sanjiv R. Das, Santa Clara University, srdas@scu.edu
- Tarang Goyal, Santa Clara University, tgoyal@scu.edu
- Mohini Yadav, Santa Clara University, myadav@scu.edu

Abstract

How capable are time-series foundation models? Using Chronos-2, an open-source time-series foundation model, we evaluate pretrained time-series models for economic and financial forecasting with an emphasis on whether multivariate (MV) inputs improve accuracy relative to univariate (UV) baselines. The study covers two panels---the Magnificent-7 equities and U.S. Treasury interest rates---as well as a combined panel, using rolling monthly evaluations from 2000--2025. We vary input window lengths and forecast horizons and report RMSE and MAPE. Across datasets, MV forecasts consistently outperform UV forecasts, with especially strong gains for interest rates and meaningful improvements for equities. Series-level comparisons show MV improvements in every case, and error dispersion is generally lower under MV inputs. We also provide parameter-heatmap and time-series visualizations. However, mixing time series across equity and interest rate markets reduces forecast accuracy, indicating that adding noisy context degrades model performance. Overall, the results indicate that foundation models can leverage cross-series information to improve forecast accuracy in finance, and that the benefits are strongest when related series are modeled jointly under disciplined rolling protocols.

Can Time-Series Foundation Models Generate Improved Multivariate Forecasts?

April 29, 2026

Abstract

How capable are time series foundation models? Using Chronos-2, an open-source time-series foundation model, we evaluate pretrained time-series models for economic and financial forecasting with an emphasis on whether multivariate (MV) inputs improve accuracy relative to univariate (UV) baselines. Just as LLMs generate sequences of words, Chronos-2 is trained on thousands of time series, and generates forward sequences of data based on multivariate histories. The study covers two panels—the Magnificent-7 equities and U.S. Treasury interest rates—as well as a combined panel, using rolling monthly evaluations from 2000–2025. We vary input window lengths and forecast horizons and report RMSE and MAPE. Across datasets, MV forecasts consistently outperform UV forecasts, with especially strong gains for interest rates and meaningful improvements for equities. Series-level comparisons show MV improvements in every case, and error dispersion is generally lower under MV inputs. We also provide parameter-heatmap and time-series visualizations. However, mixing time series across equity and interest rate markets reduces forecast accuracy, indicating that adding noisy context degrades model performance. Overall, the results indicate that foundation models can leverage cross-series information to improve forecast accuracy in finance, and that the benefits are strongest when related series are modeled jointly under disciplined rolling protocols.

1 Introduction

Time-series forecasting in finance remains a contested area for machine learning. While deep learning has transformed NLP and computer vision, its advantage in financial forecasting is less clear. Recent studies of U.S. Treasury yield curve forecasting report that classical econometric baselines often outperform modern machine-learning models across long spans of daily data, with few learning approaches delivering competitive gains. These findings underscore an open question: when do multivariate methods provide reliable improvements over univariate forecasting? Expanding data from univariate to multivariate prompts the standard concern of whether

it introduces signal or noise.

The emergence of foundation models—large, self-supervised models adapted to downstream tasks—has begun to reshape forecasting. These models promise both scale and transfer, exploiting shared structure across diverse series to improve accuracy in noisy or data-limited domains. Yet evidence from competitions and applied finance (Singh et al., 2025) suggests that pure deep-learning solutions often struggle without hybridization or careful engineering.

This paper evaluates if foundation-model forecasting delivers practical improvements in finance, focusing on Chronos-2, an open-source model (GitHub: <https://github.com/amazon-science/chronos-forecasting>). Just as large language models generate a sequence of words, Chronos-2 is trained on thousands of time series, and generates forward sequences of data based on multivariate histories. We hypothesize that properly chosen multivariate (MV) inputs can help models extract cross-series information and outperform univariate (UV) baselines. We test this across two panels—the Magnificent-7 equities and U.S. Treasury rates—and a combined panel to assess if joint modeling improves results.

The contribution of the paper is threefold. First, we provide a systematic comparison of MV versus UV forecasting across equities and rates with a consistent rolling evaluation protocol. We find that the MV approach beats the UV one. Second, we examine whether MV gains differ by asset class and whether joint stock and interest rates modeling improves accuracy. In this case, we find that adding a different class of time series does not improve forecasting, highlighting that expanding the basis set for modeling does not always improve forecast accuracy. Indeed, there is a signal to noise tradeoff that applies here: additional context is not always useful. Third, we offer large-scale outputs—error metrics, forecasts with uncertainty quantiles, and visualization-ready artifacts—to support robust evaluation and downstream use. Collectively, these results clarify when foundation models add value to financial forecasting.

1.1 Literature Review

The literature on financial time-series forecasting spans classical econometrics, machine learning (ML), and deep learning. Standard statistical tools remain central and are documented in foundational texts (Shumway and Stoffer, 2025). In fixed-income, yield curve forecasting has been explored via no-arbitrage models (Moench, 2006) and forecast combinations (Caldeira et al., 2013). ML methods show mixed results, with recent studies indicating that classical econometric baselines often outperform ML and deep-learning models across daily data (Singh et al., 2025; Sambasivan and Das, 2017). Research on bond risk premia further highlights the potential and limitations of ML for term-structure prediction (Bianchi et al., 2020; Hoogteijling et al., 2021).

For equities, surveys emphasize deep learning’s broad adoption but inconsistent performance across tasks (Sezer et al., 2019). Comparisons between ARIMA, LSTM, and Prophet show that model choice is highly data-dependent, with traditional methods remaining competitive (Sunki et al., 2024). Experts argue for rigorous validation and robust benchmarks to avoid overfitting and spurious gains (López de Prado, 2018). These themes motivate the careful use of UV baselines and out-of-sample evaluation in this study.

Recent deep-learning advances include specialized architectures like NHITS (Challu et al., 2023) and time-series foundation models. Chronos is a large, self-supervised model designed for cross-domain transfer (Ansari et al., 2024). This paper tests whether MV inputs systematically improve accuracy for equities and rates using consistent rolling protocols and multiple metrics. We specifically focus on the newest generation model, Chronos-2 (Ansari et al., 2025).

2 Methodology

This section outlines our research questions and the Chronos-2 forecasting pipeline, including data scope, rolling evaluation, and analysis artifacts.

2.1 Research Questions

We address three questions. First, does multivariate (MV) forecasting outperform univariate (UV) methods within the Chronos-2 framework? Second, do MV gains vary between asset classes (stocks vs. interest rates)? Third, does joint modeling of stocks and rates improve accuracy over separate modeling?

2.2 Overview of the Chronos-2 model

Chronos-2 advances time-series foundation models by enabling zero-shot, multivariate, and covariate-informed forecasting. It moves beyond single-series histories to joint inference across related series without task-specific retraining, simplifying deployment and improving transfer to new domains. Architecturally, Chronos-2 uses a group-attention mechanism for in-context learning across series sets. A “group” can include multivariate variates, related targets, or targets with covariates. This design facilitates information sharing and cross-series structure learning at inference time. To build these capabilities, the model is pretrained on synthetic multivariate constructions. Chronos-2 achieves state-of-the-art results on benchmarks like GIFT-Eval and Chronos Benchmark II, particularly for covariate-conditioned tasks. The methodology transforms forecasting into a universal task via a transformer architecture and innovative pre-training (see Figure 1 in the Chronos-2 technical report, (Ansari et al., 2025)). Key components include:

1. **Tokenization and Robust Scaling.** Data is normalized via robust scaling and a $\sinh^{-1}$ transformation to stabilize variance. Categorical covariates are encoded, and data is augmented with meta-features like a relative time index and a binary mask for missing values.
2. **Patch-Based Representation.** Chronos-2 divides the series and meta-features into non-overlapping patches of length P . These are mapped into embeddings via a residual network, with a REG token serving as an attention sink between history and future patches.
3. **Core Architecture: Time and Group Attention.** The encoder-only transformer alternates between: (a) Time Attention, which aggregates temporal information using RoPE, and (b) Group Attention, the core innovation allowing cross-series information sharing within a group at each patch index.
4. **Multivariate and Covariate Integration.** Tasks are managed via Group IDs. In multivariate tasks, related variates share an ID to model dependencies. Known future covariates are included in the future input matrix W , allowing for conditioned forecasts.
5. **Training and Quantile Forecasting.** Trained in two stages (context lengths up to 8192), the model uses quantile regression to produce probabilistic forecasts across 21 quantiles. Its in-context learning capabilities are largely derived from synthetic multivariate datasets.

In sum, the methodology of Chronos-2 involves an encoder-only transformer architecture that utilizes a

novel group attention mechanism and robust scaling to perform zero-shot univariate, multivariate, and covariate-informed forecasting by sharing information across related time series in-context.

2.3 Implementation

The data consist of three panels: the Magnificent-7 stocks (AAPL, AMZN, GOOGL, MSFT, NFLX, NVDA, TSLA) ($K = 7$), Treasury interest rates for ten different maturities ($K = 10$), and a combined panel of all series ($K = 17$). This design supports both within-group and cross-group forecasting comparisons.

We use input window lengths $n \in \{126, 252, 504, 756\}$ trading days, and forecast horizons $m \in \{21, 63\}$ trading days to cover short to medium histories (0.5 to 3 years) and near- to mid-term prediction horizons (1 to 3 months).

The evaluation period spans 2000–2025 with monthly rolling forecasts to reflect realistic re-estimation and deployment cadence. Forecast accuracy is measured using RMSE (root mean squared error) and MAPE (mean absolute percentage error) to capture both scale-dependent and scale-free error behavior.

To assess temporal robustness, we compare performance across pre-2023 and post-2023 regimes using a training cutoff that separates the sample. We will also argue that this is a check on information leakage in the model.

3 Results

3.1 Forecasts within asset classes

This paper examines improvements when using multivariate (MV) instead of univariate (UV) forecasting. Table 1 summarizes performance by dataset and mode across all tickers and period combinations. MV forecasts outperform UV for both rates and stocks, showing significant MAPE and RMSE reductions in rates and moderate reductions in stocks. Forecast error standard deviation is also lower for MV. Overall, multivariate models in Chronos-2 provide robust, easily implemented forecasts.

Table 1: Average performance by dataset and mode. Lower mean error across all experiments (for values of input length $n = \{126, 252, 504, 756\}$, and forecast period $m = \{21, 63\}$) measured by MAPE and RMSE is better. Data for the rolling periods is from 2000 through 2025. Predictions periods begin 3 years after the start of the dataset.

Dataset	Mode	MAPE		RMSE	
		mean	std	mean	std
Rates	MV	0.0497	0.1673	0.0418	0.0445
Rates	UV	0.1233	0.2213	0.1865	0.1783
Stocks	MV	0.0706	0.0728	3.9619	9.0440
Stocks	UV	0.0844	0.0834	5.0395	11.5710

Table 2 reports UV vs. MV results by series. MV improves MAPE and RMSE for every rate and stock. This clear win suggests that additional related time series add value when used within a pre-trained foundation time series model.

Figure 1 breaks down errors by input length n and forecast length m . MV consistently yields smaller errors than UV. Accuracy predictably decreases for the 63-day horizon compared to 21 days, but input length n (from six months to three years) has little impact on results.

MAPE error time series are shown in Figure 2. UV errors are consistently higher, especially for interest rates. Errors appear to rise at higher asset levels despite using percentage metrics. Overall, Figures 1 and 2 confirm that MV forecasts are more accurate.

3.2 Forecasts with combined asset classes

This section assesses a "small world" model combining rates and stocks to see if cross-domain data improves results. We test whether forecasting a series using both families of data exploits additional cross-market information.

Using the period July 2010 to December 2025 for consistency, Table 3 shows that cross-market data does not improve forecasts and may slightly degrade quality by adding noise. In foundation models, this unrelated context may distract the model from the target domain.

3.3 Leakage

We address potential "leakage"—the risk that test data appeared in pre-training. Chronos models exclude the financial series used here. Chronos-2 (Ansari et al., 2025) focuses on training diversity, using synthetic data via TSI (Bahrpeyma et al., 2021) and TCM (Runge et al., 2023) approaches. While the original model (Ansari et al., 2024) utilized some macroeconomic (Godaheva et al., 2021) and financial datasets (Makridakis et al.,

Table 2: UV vs. MV comparison by series (lower is better). Improvements are mean UV error minus mean MV error. Results are averaged across values of input length $n = \{126, 252, 504, 756\}$, and forecast period $m = \{21, 63\}$. Data for the rolling periods is from 2000 through 2025. Predictions periods begin 3 years after the start of the dataset.

Dataset	Series	MAPE (MV)	MAPE (UV)	RMSE (MV)	RMSE (UV)	MAPE Imp.	RMSE Imp.
Rates	DGS3MO	0.2355	0.3186	0.0618	0.1227	0.0831	0.0609
Rates	DGS6MO	0.0833	0.1770	0.0455	0.1257	0.0937	0.0802
Rates	DGS1	0.0572	0.1326	0.0454	0.1424	0.0754	0.0970
Rates	DGS2	0.0334	0.1270	0.0413	0.1845	0.0936	0.1432
Rates	DGS3	0.0243	0.1181	0.0388	0.2068	0.0938	0.1680
Rates	DGS5	0.0171	0.0972	0.0359	0.2232	0.0801	0.1873
Rates	DGS7	0.0134	0.0827	0.0345	0.2278	0.0693	0.1933
Rates	DGS10	0.0114	0.0703	0.0344	0.2184	0.0589	0.1840
Rates	DGS20	0.0099	0.0572	0.0364	0.2120	0.0473	0.1756
Rates	DGS30	0.0113	0.0520	0.0442	0.2012	0.0407	0.1570
Stocks	AAPL	0.0599	0.0720	3.3016	4.1563	0.0121	0.8547
Stocks	AMZN	0.0567	0.0707	3.0825	4.0517	0.0140	0.9692
Stocks	GOOGL	0.0478	0.0635	3.1803	4.0070	0.0157	0.8267
Stocks	MSFT	0.0368	0.0466	4.2003	6.1519	0.0098	1.9516
Stocks	NFLX	0.1040	0.1145	1.7871	2.0192	0.0105	0.2321
Stocks	NVDA	0.0906	0.1107	1.5994	1.8420	0.0201	0.2426
Stocks	TSLA	0.1098	0.1242	13.6796	16.7870	0.0144	3.1074

2020), they did not include our specific series, especially post-2023 data.

Two points counter leakage: (i) if leakage existed, MV and UV forecasts would show little difference, yet significant gains appear in Tables 1 through 3; (ii) post-2023 forecasts—which could not be in training—are actually more accurate (Figure 3), especially for rates. These results suggest leakage is not a factor and that foundation models excel when using multivariate series within the same contextual domain.

4 Concluding Discussion

This paper evaluates Chronos-2 as a foundation-model for financial forecasting, specifically testing if multivariate (MV) inputs improve accuracy over univariate (UV) baselines. Across equities and U.S. Treasury rates, MV forecasts consistently outperform UV counterparts, showing significant MAPE and RMSE reductions. These results demonstrate that foundation models effectively capitalize on cross-series information when inputs are properly structured. These results are some of the first that support this class of foundation models for use in financial forecasting.

Our findings contribute to the debate on machine learning in finance. Consistent with evidence that classical benchmarks are resilient, our results show that gains depend on leveraging multivariate structures within disciplined rolling protocols. The primary advantage of foundation models likely stems from their ability to transfer information across related series and accommo-

date numerous inputs—a task often difficult for traditional econometric models.

Several limitations remain. This study focuses on one foundation model and a specific asset set. The evaluation uses fixed windows and horizons; different data frequencies or rolling cadences might yield alternative results. Furthermore, while we report standard error metrics, specific use cases may require deeper analysis of tail risk, directional accuracy, and economic utility.

Future research can extend this framework in three directions. First, expanding to a larger "world" panel—including macroeconomics, commodities, and international data—could test transfer learning at scale. Second, integrating uncertainty calibration, such as conformal prediction, would enhance forecast interval reliability. Finally, hybrid designs merging econometric baselines with foundation-model embeddings may provide robust performance alongside interpretability.

Table 3: Combined dataset forecasts. UV vs. MV comparison by series (lower is better). Improvements are mean UV error minus mean MV error. Results are averaged across values of input length $n = \{126, 252, 504, 756\}$, and forecast period $m = \{21, 63\}$. Data is for the period July 2010 through December 2025, the period when there is complete data for all stocks and rates. For comparison, we show the results when forecasts are based on the individual datasets for Rates and Stocks, as well as when forecasts are based on the combination of both datasets. Predictions periods begin 3 years after the start of the dataset. (As expected the UV forecast errors do not change.)

Dataset	Series	MAPE (MV)	MAPE (UV)	RMSE (MV)	RMSE (UV)	MAPE Imp.	RMSE Imp.
Rates	DGS3MO	0.2304	0.2849	0.0597	0.1052	0.0545	0.0455
Rates	DGS6MO	0.0937	0.1714	0.0500	0.1179	0.0777	0.0679
Rates	DGS1	0.0609	0.1289	0.0475	0.1363	0.0680	0.0888
Rates	DGS2	0.0366	0.1216	0.0426	0.1727	0.0850	0.1301
Rates	DGS3	0.0257	0.1111	0.0378	0.1883	0.0854	0.1505
Rates	DGS5	0.0185	0.0959	0.0341	0.2025	0.0774	0.1684
Rates	DGS7	0.0144	0.0843	0.0321	0.2086	0.0699	0.1765
Rates	DGS10	0.0128	0.0759	0.0325	0.2045	0.0631	0.1720
Rates	DGS20	0.0116	0.0628	0.0353	0.1993	0.0512	0.1640
Rates	DGS30	0.0122	0.0573	0.0396	0.1923	0.0451	0.1527
Stocks	AAPL	0.0500	0.0621	5.6015	7.1124	0.0121	1.5109
Stocks	AMZN	0.0462	0.0617	5.2798	7.0294	0.0155	1.7496
Stocks	GOOGL	0.0411	0.0545	4.8211	6.0685	0.0134	1.2474
Stocks	MSFT	0.0305	0.0453	6.6967	10.1334	0.0148	3.4367
Stocks	NFLX	0.0792	0.0902	3.0946	3.4715	0.0110	0.3769
Stocks	NVDA	0.0743	0.0980	2.7598	3.3253	0.0237	0.5655
Stocks	TSLA	0.1090	0.1258	16.6789	20.7258	0.0168	4.0469
Dataset	Series	MAPE (MV)	MAPE (UV)	RMSE (MV)	RMSE (UV)	MAPE Imp.	RMSE Imp.
Combined	DGS3MO	0.2392	0.2849	0.0655	0.1052	0.0457	0.0397
Combined	DGS6MO	0.1002	0.1714	0.0543	0.1179	0.0712	0.0636
Combined	DGS1	0.0671	0.1289	0.0518	0.1363	0.0618	0.0845
Combined	DGS2	0.0392	0.1216	0.0440	0.1727	0.0824	0.1287
Combined	DGS3	0.0256	0.1111	0.0374	0.1883	0.0855	0.1509
Combined	DGS5	0.0189	0.0959	0.0346	0.2025	0.0770	0.1679
Combined	DGS7	0.0145	0.0843	0.0329	0.2086	0.0698	0.1757
Combined	DGS10	0.0126	0.0759	0.0327	0.2045	0.0633	0.1718
Combined	DGS20	0.0121	0.0628	0.0367	0.1993	0.0507	0.1626
Combined	DGS30	0.0129	0.0573	0.0414	0.1923	0.0444	0.1509
Combined	AAPL	0.0519	0.0621	5.6524	7.1124	0.0102	1.4600
Combined	AMZN	0.0470	0.0617	5.1538	7.0294	0.0147	1.8756
Combined	GOOGL	0.0399	0.0545	4.5781	6.0685	0.0146	1.4904
Combined	MSFT	0.0318	0.0453	6.7621	10.1334	0.0135	3.3713
Combined	NFLX	0.0786	0.0902	3.0877	3.4715	0.0116	0.3838
Combined	NVDA	0.0761	0.0980	2.8012	3.3253	0.0219	0.5241
Combined	TSLA	0.1114	0.1258	16.8843	20.7258	0.0144	3.8415




Figure 1: Parameter heatmap comparing MV and UV forecasting accuracy across window lengths and horizons (based on MAPE). The upper plot is for Rates and the lower one is for Stocks. In each plot, the left side is for the 21 working day forecast (1 month forecast) and the right side is for the 63 working day forecast (3 month forecast)




Figure 2: Time series of MV and UV MAPE errors over the evaluation period. This is for the dataset running from July 2010 through December 2025.




Figure 3: Comparison of forecast accuracy pre-2023 and post-2023. Data used is for the period July 2010 through December 2025.

References

- Ansari, A. F., O. Shchur, J. Küken, A. Auer, B. Han, P. Mercado, S. S. Rangapuram, H. Shen, L. Stella, X. Zhang, M. Goswami, S. Kapoor, D. C. Maddix, P. Guerron, T. Hu, J. Yin, N. Erickson, P. M. Desai, H. Wang, H. Rangwala, G. Karypis, Y. Wang, and M. Bohlke-Schneider (2025, October). Chronos-2: From Univariate to Universal Forecasting. [arXiv:2510.15821 \[cs\]](#).
- Ansari, A. F., L. Stella, C. Turkmen, X. Zhang, P. Mercado, H. Shen, O. Shchur, S. S. Rangapuram, S. P. Arango, S. Kapoor, J. Zschiegner, D. C. Maddix, H. Wang, M. W. Mahoney, K. Torkkola, A. G. Wilson, M. Bohlke-Schneider, and Y. Wang (2024). Chronos: Learning the Language of Time Series. [arXiv: 2403.07815](#).
- Bahrpeyma, F., M. Roantree, P. Cappellari, M. Scriney, and A. McCarren (2021, January). A Methodology for Validating Diversity in Synthetic Time Series Generation. [MethodsX](#) 8, 101459.
- Bianchi, D., M. Büchner, and A. Tamoni (2020). Bond Risk Premiums with Machine Learning. [The Review of Financial Studies](#) 34(2), 1046–1089.
- Caldeira, J., G. V. Moura, and A. A. Santos (2013). Predicting the Yield Curve using Forecast Combinations. Publication Title: SSRN Electronic Journal.
- Challu, C., K. G. Olivares, B. N. Oreshkin, F. G. Ramirez, M. M. Canseco, and A. Dubrawski (2023, June). NHITS: Neural Hierarchical Interpolation for Time Series Forecasting. [Proceedings of the AAAI Conference on Artificial Intelligence](#) 37(6), 6989–6997.
- Godahehwa, R., C. Bergmeir, G. I. Webb, R. J. Hyndman, and P. Montero-Manso (2021, May). Monash Time Series Forecasting Archive. [arXiv:2105.06643 \[cs\]](#).
- Hoogteijling, T., M. P. E. Martens, and M. v. d. Wel (2021, April). Forecasting Bond Risk Premia using Stationary Yield Factors.
- López de Prado, M. (2018, January). [Advances in Financial Machine Learning](#). Wiley.
- Makridakis, S., E. Spiliotis, and V. Assimakopoulos (2020). The M4 Competition: 100,000 time-series and 61 Forecasting Methods. [International Journal of Forecasting](#) 36(1), 54–74.
- Moench, E. (2006). Forecasting the Yield Curve in a Data-Rich Environment: A No-Arbitrage Factor-Augmented VAR Approach. [SSRN Electronic Journal](#).
- Runge, J., A. Gerhardus, G. Varando, V. Eyring, and G. Camps-Valls (2023, July). Causal inference for time series. [Nature Reviews Earth and Environment](#) 4, 487–505. ADS Bibcode: 2023NRvEE...4..487R.
- Sambasivan, R. and S. Das (2017). A Statistical Machine Learning Approach to Yield Curve Forecasting. In [2017 International Conference on Computational Intelligence in Data Science \(ICCIDS\)](#).
- Sezer, O. B., M. U. Gudelek, and A. M. Ozbayoglu (2019, November). Financial Time Series Forecasting with Deep Learning : A Systematic Literature Review: 2005-2019. [arXiv:1911.13288 \[cs\]](#).
- Shumway, R. H. and D. S. Stoffer (2025). [Time Series Analysis and Its Applications: With R Examples](#). Springer Texts in Statistics. Cham: Springer Nature Switzerland.
- Singh, A., T. Ogunfunmi, and S. Das (2025). Yield Curve Forecasting using Machine Learning and Econometrics: A Comparative Analysis. [Journal of Investment Management](#) 23(4), 37–58.
- Sunki, A., C. SatyaKumar, G. Surya Narayana, V. Koppera, and M. Hakeem (2024). Time series forecasting of stock market using ARIMA, LSTM and FB prophet. [MATEC Web of Conferences](#) 392, 01163.

A Appendix: AI-Driven Workflow to Generate the Paper

This appendix outlines how the paper was produced with the assistance of Artificial Intelligence.

Declaration of generative AI and AI-assisted technologies in the manuscript preparation process

During the preparation of this work the author(s) used many AI tools in order to develop this paper. After using these tools/services, the author(s) reviewed and edited the content as needed and take(s) full responsibility for the content of the published article.

Detailed Workflow Step-by-Step

Each step below was followed by human review, if AI was involved.

1. **Conceptualization:** Students are provided an initial outline.
2. **Initial Coding:** Students prepare a Google Colab notebook with assistance from Gemini.
3. **Refinement:** The professor updates the notebook, modularizing the code and updating logic, also utilizing Gemini for coding assistance.
4. **Data Preparation:** Results are finalized in Colab, consisting of printed dataframes, plots, and textual outputs.
5. **Environment Setup:** A $\text{\LaTeX}$ shell for the paper is established in OpenAI Prism (<https://openai.com/prism/>).
6. **Outlining:** The structure is set up with blank sections: Introduction, Methodology, Results, and Concluding Discussion (including relevant subsections).
7. **Table Conversion:** Printed dataframes from Colab are fed into Gemini with prompts to convert them into $\text{\LaTeX}$ tables. These are then integrated into the Results section.
8. **Visuals:** Figures are added to the project folder, and the corresponding $\text{\LaTeX}$ code for figure inclusion is written.
9. **Drafting Results:** Prism is tasked with drafting text for the Results section to explain the tables and figures. Notably, the underlying model (GPT-5.2) inferred results directly from the tables to generate the write-up.
10. **Methodology:** The Chronos-1 and Chronos-2 papers, along with experiment descriptions from the Colab notebook, are uploaded to Gemini. Gemini generates the Methodology section, which is then moved to the draft.
11. **Specific Subsections:** A subsection on data leakage is added; the reasoning is provided by the authors while the LLM handles the prose.
12. **Literature Review (Part 1):** A Bib $\text{\TeX}$ file from Zotero is uploaded. Prism is prompted to generate a literature review based on the citations as they relate to the Methodology and Results.
13. **Introduction and Literature Review (Part 2):** Previously published related papers are uploaded. Prism generates the Introduction and updates the literature review for context.
14. **Conclusions:** Using the full paper as context, Prism generates the Concluding Discussion.
15. **Abstract:** Prism is prompted to synthesize an abstract for the completed draft.
16. **Human Editing:** Final human-led editing passes are performed to improve narrative flow, reformat the $\text{\LaTeX}$ layout, and ensure technical accuracy.

Self Review and Revision Process

Gemini was also utilized to generate a mock referee report (`review.pdf`). This report informed a response document (`responses.md`) created with the following prompt:

“Please read the file ‘review.pdf’ which is a referee report on the paper ‘chronos2.pdf’ – then suggest changes and improvements to the paper in a document titled ‘response.md’ that outlines the changes in response to cited sections of ‘review.pdf’ as a response note to the review.”

Final edits were made to the manuscript based on these AI-suggested improvements.